

Vignesh Girish

Senior Data Platform Engineer

+44-7760973695 | vignesh.malhotra15@gmail.com | London, United Kingdom
[linkedin.com/in/vignesh-malhotra](https://www.linkedin.com/in/vignesh-malhotra) | github.com/Nerdy-Thanos

SUMMARY

Data platform engineer with 6+ years of experience building high-throughput batch and streaming pipelines and self-service cloud platforms on AWS, including 3+ years at tier-1 UK banks (Citi, Barclays). Delivered near real-time systems processing millions of records daily, along with multi-tenant platform tooling — Terraform modules, Helm charts, CI/CD pipelines — adopted by engineering teams across the organisation. Owns systems end to end: architecture, infrastructure-as-code, deployment and production support.

WORK EXPERIENCE

Citi UK

March 2026 – Present

Cloud Platform Engineer

- Enabled Databricks Classic and Serverless workspaces for tenant teams via shared Terraform modules, enforcing OPA policies and organisational guardrails through a plug-and-play deployment model.
- Cut CI/CD test pipeline runtimes from 35 to 10 minutes by orchestrating ephemeral vCluster virtual clusters on EKS and GKE for automated API testing.
- Built shared Helm charts for running optimised, large-scale Spark jobs on EKS and GKE, consumed by tenant teams through private APIs.
- Migrated 100+ repositories from Bitbucket to GitHub Enterprise and re-platformed CI/CD from Jenkins to Tekton and Harness as part of an organisation-wide modernisation.

Barclays UK

July 2024 – February 2026

Senior AWS Data Platform Engineer

- Designed and deployed an end-to-end pub-sub notification service providing a standard publish and consume interface for applications across the bank.
- Led development and deployment of a data pipeline processing ~400,000 nested-JSON API records daily using AWS Lambda, S3 and Glue.
- Implemented a scalable transformation layer using dbt-core on AWS Fargate, managing 40+ models delivering clean, query-optimised data to Apache Iceberg tables.
- Orchestrated near real-time streaming workflows with AWS Step Functions and DynamoDB, achieving 2–5 minute end-to-end latency and 99.9% uptime.
- Created reusable CloudFormation templates shared across data domains, cutting deployment time by 40% and preventing resource conflicts between stacks.

Barclays UK

February 2023 – June 2024

Data Platform Developer

- Reduced infrastructure provisioning time by 30% using CloudFormation and Jenkins to deploy Glue jobs, S3 buckets, RDS instances and Lambda functions.
- Accelerated data ingestion with Scala Spark scripts on AWS Glue and Kafka, processing 40 source tables into Aurora PostgreSQL for analytics.
- Built automated CI/CD pipelines with Jenkins and CloudFormation, cutting manual integration and deployment time by 75%.
- Developed a distributed serverless ETL pipeline using AWS Step Functions, processing 2M+ rows in 4–5 minutes with outputs stored in Aurora PostgreSQL.

Adsum Technologies

September 2022 – January 2023

Data Engineer

- Built automated Apache Airflow pipelines integrating Xero accounting data, refreshing daily financials and processing 10,000+ transactions per batch into Snowflake.
- Deployed a Kubernetes cluster on AWS EKS to manage containerised Airflow workflows, improving scalability and reducing pipeline downtime by 25%.
- Integrated Kafka to stream financial transaction logs from open-banking connections, processing 15,000+ events daily with near real-time updates to Snowflake.

- Built a lead-generation microservice extracting financial data from Companies House PDF filings via OCR into RDS PostgreSQL, saving 40–50 hours of market research per week.

CavinKare

June 2019 – August 2021

Data Engineer

- Developed a deep learning face recognition system in Python, scaled to 500 employees, replacing traditional biometrics with 95% accuracy.
- Designed a computer-vision safety monitoring pipeline on Azure ML Studio with 92% accuracy, reducing off-time security hours by 30%.
- Set up distributed Spark pipelines on Azure Databricks, later migrated to AWS Glue for scalability, processing 500GB+ of sales data in 1–2 minutes.

SKILLS

- **Languages** – Python, Scala, Java, SQL
- **Infrastructure as Code** – Terraform, CloudFormation, Helm, AWS SAM
- **AWS** – S3, Lambda, Glue, Step Functions, DynamoDB, EKS, ECS, Fargate, Redshift, Aurora / RDS PostgreSQL, IAM, KMS, VPC, CloudWatch
- **Other Clouds** – GCP (GKE, BigQuery), Azure (Databricks, ML Studio)
- **Data & Streaming** – Spark, Kafka, dbt, Apache Iceberg, Snowflake, Databricks, Apache Airflow, batch and streaming ETL/ELT
- **Containers & Orchestration** – Docker, Kubernetes (EKS, GKE, vCluster)
- **CI/CD & Source Control** – Jenkins, GitLab CI, GitHub Actions, Tekton, Harness, AWS CodeDeploy, OPA policies, Git, GitHub Enterprise, Bitbucket

EDUCATION

University of Glasgow

September 2021 – December 2022

MSc Data Science

- MSc project: 2.5D image generation using deep learning.

CERTIFICATIONS

- AWS Certified Solutions Architect – Associate
- Databricks Data Engineer Associate